

LEARNING THROUGH DIALOGUE INTERACTIONS

Jiwei Li, Alexander H. Miller, Sumit Chopra, Marc’Aurelio Ranzato, Jason Weston

Facebook AI Research,

New York, USA

{jiwel, ahm, spchopra, ranzato, jase}@fb.com

ABSTRACT

A good dialogue agent should have the ability to interact with users. In this work, we explore this direction by designing a simulator and a set of synthetic tasks in the movie domain that allow the learner to interact with a teacher by both asking and answering questions. We investigate how a learner can benefit from asking questions in both an offline and online reinforcement learning setting. We demonstrate that the learner improves when asking questions. Our work represents a first step in developing end-to-end learned interactive dialogue agents.

1 INTRODUCTION

When a student is asked a question by a teacher, but is not confident about the answer, they may ask for clarification or hints. A good conversational agent (a learner/bot/student) should have this ability to interact with a dialogue partner (the teacher/user). However, recent efforts have mostly focused on learning through fixed answers provided in the training set, rather than through interactions. In that case, when a learner encounters a confusing situation such as an unknown surface form (phrase or structure), a semantically complicated sentence or an unknown word, the agent will either make a (usually poor) guess or will redirect the user to other resources (e.g., a search engine, as in Siri). Humans, in contrast, can adapt to many situations by asking questions.

We identify three categories of mistakes a learner can make during dialogue¹: (1) the learner has problems understanding the surface form of the text of the dialogue partner, e.g., the phrasing of a question; (2) the learner has a problem with reasoning, e.g. they fail to retrieve and connect the relevant knowledge to the question at hand; (3) the learner lacks the knowledge necessary to answer the question in the first place – that is, the knowledge sources the student has access to do not contain the needed information.

All the situations above can be potentially addressed through interaction with the dialogue partner. Such interactions can be used to *learn to perform better in future dialogues*. If a human student has problems understanding a teacher’s question, they might ask the teacher to clarify the question. If the student doesn’t know where to start, they might ask the teacher to point out which known facts are most relevant. If the student doesn’t know the information needed at all, they might ask the teacher to tell them the knowledge they’re missing, writing it down for future use.

In this work, we try to bridge the gap between how a human and an end-to-end machine learning dialogue agent deal with these situations: our student has to *learn how to learn*. We hence design a simulator and a set of synthetic tasks in the movie question answering domain that allow a bot to interact with a teacher to address the issues described above. Using this framework, we explore how a bot can benefit from interaction by asking questions in both offline supervised settings and online reinforcement learning settings, as well as how to choose when to ask questions in the latter setting. In both cases, we find that the learning system improves through interacting with users.

2 RELATED WORK

Learning language through interaction and feedback can be traced back to the 1950s, when Wittgenstein argued that the meaning of words is best understood from their use within given language

¹This list is not exhaustive; for example, we do not address a failure in the dialogue generation stage.

games (Wittgenstein, 2010). The direction of interactive language learning through language games has been explored in the early seminal work of Winograd (Winograd, 1972), and in the recent SHRD-LURN system (Wang et al., 2016). In a broader context, the usefulness of feedback and interactions has been validated in the setting of multiple language learning, such as second language learning (Bassiri, 2011) and learning by students (Higgins et al., 2002; Latham, 1997; Werts et al., 1995).

In the context of dialogue, with the recent popularity of deep learning models, many neural dialogue systems have been proposed. These include the *chit-chat* type *end-to-end dialogue systems* (Vinyals & Le, 2015; Li et al., 2015; Sordani et al., 2015), which directly generate a response given the previous history of user utterance. It also include a collection of *goal-oriented* dialogue systems (Wen et al., 2016; Su et al., 2016; Bordes & Weston, 2016), which complete a certain task such as booking a ticket or making a reservation at a restaurant. Another line of research focuses on supervised learning for question answering from dialogues (Dodge et al., 2015; Weston, 2016), using either a given database of knowledge (Bordes et al., 2015; Miller et al., 2016) or short stories (Weston et al., 2015). As far as we know, current dialogue systems mostly focus on learning through fixed supervised signals rather than interacting with users.

Our work is closely related to the recent work of Weston (2016), which explores the problem of learning through conducting conversations, where supervision is given naturally in the response during the conversation. Their work introduced multiple learning schemes from dialogue utterances. In particular the authors discussed *Imitation Learning*, where the agent tries to learn by imitating the dialogue interactions between a teacher and an expert student; *Reward-Based Imitation Learning*, which only learns by imitating the dialogue interactions which have correct answers; and *Forward Prediction*, which learns by predicting the teacher’s feedback to the student’s response. Despite the fact that Forward Prediction does not use human-labeled rewards, the authors show that it yields promising results. However, their work did not fully explore the ability of an agent to learn via questioning and interaction. Our work can be viewed as a natural extension of theirs.

3 THE TASKS

In this section we describe the dialogue tasks we designed². They are tailored for the three different situations described in Section 1 that motivate the bot to ask questions: (1) *Question Clarification*, in which the bot has problems understanding its dialogue partner’s text; (2) *Knowledge Operation*, in which the bot needs to ask for help to perform reasoning steps over an existing knowledge base; and (3) *Knowledge Acquisition*, in which the bot’s knowledge is incomplete and needs to be filled.

For our experiments we adapt the WikiMovies dataset (Weston et al., 2015), which consists of roughly 100k questions over 75k entities based on questions with answers in the open movie dataset (OMDb). The training/dev/test sets respectively contain 181638 / 9702 / 9698 examples. The accuracy metric corresponds to the percentage of times the student gives correct answers to the teacher’s questions.

Each dialogue takes place between a teacher and a bot, generated using a simulator. The bot is first presented with facts from the OMDb KB. This allows us to control the exact knowledge the bot has access to. Then, we include several teacher-bot question-answer pairs unrelated to the question the bot needs to answer, which we call *conversation histories*³. In order to explore the benefits of asking clarification questions during a conversation, for each of the three scenarios, our simulator generated data for two different settings, namely, *Question-Answering* (denoted by QA), and *Asking-Question* (denoted by AQ). For both QA and AQ, the bot needs to give an answer to the teacher’s original question at the end. The details of the simulator can be found in the appendix.

3.1 QUESTION CLARIFICATION.

In this setting, the bot does not understand the teacher’s question. We focus on a special situation where the bot does not understand the teacher because of typo/spelling mistakes, as shown in Figure 1. We intentionally misspell some words in the questions such as replacing the word “movie” with

² Code and data are available at <https://github.com/facebook/MemNN/tree/master/AskingQuestions>.

³ These history QA pairs can be viewed as distractions and are used to test the bot’s ability to separate the wheat from the chaff. For each dialogue, we incorporate 5 extra QA pairs (10 sentences).

“**movvie**” or “star” with “**sttar**”⁴. To make sure that the bot will have problems understanding the question, we guarantee that the bot has never encountered the misspellings before—the misspelling-introducing mechanisms in the training, dev and test sets are different, so the same word will be misspelled in different ways in different sets. We present two *AQ* tasks: (i) *Question Paraphrase* where the student asks the teacher to use a paraphrase that does not contain spelling mistakes to clarify the question by asking “what do you mean?”; and (ii) *Question Verification* where the student asks the teacher whether the original typo-bearing question corresponds to another question without the spelling mistakes (e.g., “Do you mean which film did Tom Hanks appear in?”). The teacher will give feedback by giving a paraphrase of the original question without spelling mistakes (e.g., “I mean which film did Tom Hanks appear in”) in *Question Paraphrase* or positive/negative feedback in *Question Verification*. Next the student will give an answer and the teacher will give positive/negative feedback depending on whether the student’s answer is correct. Positive and negative feedback are variants of “No, that’s incorrect” or “Yes, that’s right”⁵. In these tasks, the bot has access to all relevant entries in the KB.

3.2 KNOWLEDGE OPERATION

The bot has access to all the relevant knowledge (facts) but lacks the ability to perform necessary reasoning operations over them; see Figure 2. We focus on a special case where the bot will try to understand what are the relevant facts. We explore two settings: *Ask For Relevant Knowledge* (Task 3) where the bot directly asks the teacher to point out the relevant KB fact and *Knowledge Verification* (Task 4) where the bot asks whether the teacher’s question is relevant to one particular KB fact. The teacher will point out the relevant KB fact in the *Ask For Relevant Knowledge* setting or give a positive or negative response in the *Knowledge Verification* setting. Then the bot will give an answer to the teacher’s original question and the teacher will give feedback on the answer.

3.3 KNOWLEDGE ACQUISITION

For the tasks in this subsection, the bot has an incomplete KB and there are entities important to the dialogue missing from it, see Figure 3. For example, given the question “Which movie did Tom Hanks star in?”, the missing part could either be the entity that the teacher is asking about (question entity for short, which is Tom Hanks in this example), the relation entity (starred actors), the answer to the question (Forrest Gump), or the combination of the three. In all cases, the bot has little chance of giving the correct answer due to the missing knowledge. It needs to ask the teacher the answer to acquire the missing knowledge. The teacher will give the answer and then move on to other questions (captured in the conversational history). They later will come back to reask the question. At this point, the bot needs to give an answer since the entity is not new any more.

Though the correct answer has effectively been included in the earlier part of the dialogue as the answer to the bot’s question, as we will show later, many of the tasks are not as trivial as they look when the teacher reasks the question. This is because the bot’s model needs to memorize the missing entity and then construct the links between the missing entities and known ones. This is akin to the real world case where a student might make the same mistake again and again even though each time the teacher corrects them if their answer is wrong. We now detail each task in turn.

Missing Question Entity: The entity that the teacher is asking about is missing from the knowledge base. All KB facts containing the question entity will be hidden from the bot. In the example for *Task 5* in Figure 3, since the teacher’s question contains the entity Tom Hanks, the KB facts that contain Tom Hanks are hidden from the bot.

Missing Answer Entity: The answer entity to the question is unknown to the bot. All KB facts that contain the answer entity will be hidden. Hence, in *Task 6* of Figure 3, all KB facts containing the answer entity Forrest Gump will be hidden from the bot.

Missing Relation Entity: The relation type is unknown to the bot. In *Task 7* of Figure 3, all KB facts that express the relation starred actors are hidden from the bot.

⁴Many reasons could lead to the bot not understanding the teacher’s question, e.g., the teacher’s question has an unknown phrase structure, rather than unknown words. We choose to use spelling mistakes because of the ease of dataset construction.

⁵In the datasets we build, there are 6 templates for positive feedback and 6 templates for negative feedback.

<p>Task 1: Question Paraphrase (QA) kb: <u>Larry Crowne</u> directed_by <u>Tom Hanks</u> kb: <u>Forrest Gump</u> starred_actors <u>Tom Hanks</u>, Robin Wright, Gary Sinise kb: <u>Forrest Gump</u> directed_by <u>Robert Zemeckis</u> T/S : Conversation History. T : Which movie did Tom Hanks star in ? S : Larry Crowne T : Wrong. (-)</p>	<p>Task 1: Question Paraphrase (AQ): kb: <u>Larry Crowne</u> directed_by <u>Tom Hanks</u> kb: <u>Forrest Gump</u> starred_actors <u>Tom Hanks</u>, Robin Wright, Gary Sinise kb: <u>Forrest Gump</u> directed_by <u>Robert Zemeckis</u> T/S : Conversation History. T : Which movie did Tom Hanks star in ? S : What do you mean ? T : I mean which film did Tom Hanks appear in. T : Which movie did Tom Hanks star in ? S : Forrest Gump T : That's correct. (+)</p>
<p>Task 2: Question Verification (QA): kb: <u>Larry Crowne</u> directed_by <u>Tom Hanks</u> kb: <u>Forrest Gump</u> starred_actors <u>Tom Hanks</u>, Robin Wright, Gary Sinise kb: <u>Forrest Gump</u> directed_by <u>Robert Zemeckis</u> T/S : Conversation History. T : Which movie did Tom Hanks star in ? S : Forrest Gump T : That's correct. (+)</p>	<p>Task 2: Question Verification (AQ): kb: <u>Larry Crowne</u> directed_by <u>Tom Hanks</u> kb: <u>Forrest Gump</u> starred_actors <u>Tom Hanks</u>, Robin Wright, Gary Sinise kb: <u>Forrest Gump</u> directed_by <u>Robert Zemeckis</u> T/S : Conversation History. T : Which movie did Tom Hanks star in ? S : Do you mean which film did Tom Hanks appear in ? T : That's correct. (+) T : Which movie did Tom Hanks star in ? S : Larry Crowne T : Sorry, that's not right. (+)</p>

Figure 1: Sample dialogues for the two *Question Clarification* tasks (rows) using both the traditional *QA* setting (left column) and *AQ* setting (right column). In each case the same example is given for simplicity. Black text prefixed by “kb:” denotes KB knowledge that the student has access to. Blue text is spoken by the teacher, denoted by T. (+) denotes a reward of 1 (and 0 otherwise) that the teacher assigns to the bot. Red text denotes responses or questions posed by the bot, denoted by S. Brown denotes typos deliberately introduced by the authors. For the *Question Verification* setting, the student can either ask a correct (pertinent) question (as in this example) or an incorrect (irrelevant) one. The teacher will give positive or negative feedback based on the correctness of the student’s question. In our offline supervised learning experiments, the probability of asking pertinent questions and correctly answering the original question from the teacher is set to 0.5. Finally, T/S denotes 5 pairs of questions and answers that are irrelevant to the rest of the conversation.

<p>Task 3: Ask For Relevant Knowledge (AQ): kb: <u>Larry Crowne</u> directed_by <u>Tom Hanks</u> kb: <u>Forrest Gump</u> starred_actors <u>Tom Hanks</u>, Robin Wright, Gary Sinise kb: <u>Forrest Gump</u> directed_by <u>Robert Zemeckis</u> T/S : Conversation History. T : Which movie did Tom Hanks star in ? S : Can you give me a hint ? T : It is related to the kb fact that <u>Forrest Gump</u> starred_actors <u>Tom Hanks</u> <u>Robin Wright</u>, <u>Gary Sinise</u> T : Which movie did Tom Hanks star in ? S : Forrest Gump T : That's correct. (+)</p>	<p>Task 4: Knowledge Verification (AQ): kb: <u>Larry Crowne</u> directed_by <u>Tom Hanks</u> kb: <u>Forrest Gump</u> starred_actors <u>Tom Hanks</u>, Robin Wright, Gary Sinise kb: <u>Forrest Gump</u> directed_by <u>Robert Zemeckis</u> T/S : Conversation History. T : Which movie did Tom Hanks star in ? S : Does it have something to do with <u>Larry Crowne</u> directed_by <u>Tom Hanks</u>? T : No. That's wrong. (-) T : Which movie did Tom Hanks star in ? S : Forrest Gump T : That's correct. (+)</p>
---	---

Figure 2: Sample dialogues for *Knowledge Operation* tasks.

Missing Triples: The triple that expresses the relation between the question entity and the answer entity is hidden from the bot. In *Task 8* of Figure 3, the triple “Forrest Gump (question entity) starred_actors Tom Hanks (answer entity)” will be hidden.

Figure 3: Different Tasks for **Knowledge Acquisition**. Crossed lines correspond to entries of the KB whose retrieval is missed.

<p>Task 5: Missing Question Entity (QA) kb: Larry Crowne directed by Tom Hanks kb: Forrest Gump starred actors Tom Hanks kb: Forrest Gump starred actors Sally Field kb: Forrest Gump directed by Robert Zemeckis T/S : Conversation History. T : Which movie did Tom Hanks star in ? S : Forrest Gump T : That's correct. (+)</p>	<p>Task 5: Missing Question Entity (AQ) kb: Larry Crowne directed by Tom Hanks kb: Forrest Gump starred actors Tom Hanks, kb: Forrest Gump starred actors Sally Field kb: Forrest Gump directed by Robert Zemeckis T : Which movie did Tom Hanks star in ? S : I don't know. What's the answer? T : The answer is Forrest Gump. T/S : Conversation History. T : Which movie did Tom Hanks star in ? S : Forrest Gump T : That's correct. (+)</p>
<p>Task 6: Missing Answer Entity (AQ) kb: Larry Crowne directed by Tom Hanks kb: Forrest Gump starred actors Tom Hanks kb: Forrest Gump starred actors Sally Field kb: Forrest Gump directed by Robert Zemeckis T : Which movie did Tom Hanks star in ? S : I don't know. What's the answer? T : The answer is Forrest Gump. T/S : Conversation History. T : Which movie did Tom Hanks star in ? S : Forrest Gump T : That's correct. (+)</p>	<p>Task 7: Missing Relation Entity (AQ) kb: Larry Crowne directed by Tom Hanks kb: Forrest Gump starred actors Tom Hanks, kb: Forrest Gump starred actors Sally Field kb: Forrest Gump directed by Robert Zemeckis T : Which movie did Tom Hanks star in ? S : I don't know. What's the answer? T : The answer is Forrest Gump. T/S : Conversation History. T : Which movie did Tom Hanks star in ? S : Forrest Gump T : That's correct. (+)</p>
<p>Task 8: Missing Triple (AQ) kb: Larry Crowne directed by Tom Hanks kb: Forrest Gump starred actors Tom Hanks kb: Forrest Gump starred actors Sally Field kb: Forrest Gump directed by Robert Zemeckis T : Which movie did Tom Hanks star in ? S : I don't know. What's the answer? T : The answer is Forrest Gump. T/S : Conversation History. T : Which movie did Tom Hanks star in ? S : Forrest Gump T : That's correct. (+)</p>	<p>Task 9: Missing Everything (AQ) kb: Larry Crowne directed by Tom Hanks kb: Forrest Gump starred actors Tom Hanks, kb: Forrest Gump starred actors Sally Field kb: Forrest Gump directed by Robert Zemeckis T : Which movie did Tom Hanks star in ? S : I don't know. What's the answer? T : The answer is Forrest Gump. T/S : Conversation History. T : Which movie did Tom Hanks star in ? S : Forrest Gump T : That's correct. (+)</p>

Missing Everything: The question entity, the relation entity, the answer entity are all missing from the KB. All KB facts in *Task 9* of Figure 3 will be removed since they either contain the relation entity (i.e., starred actors), the question entity (i.e., Forrest Gump) or the answer entity Tom Hanks.

4 TRAIN/TEST REGIME

We now discuss in detail the regime we used to train and test our models. Assuming access to the simulator, our objective was twofold. First, we wanted to test the usefulness of asking questions. Second, we wanted to train our student bot to learn when to ask questions and what questions to ask. In order to accomplish these two objectives we explored training our models using two methodologies, namely, Offline Supervised Learning and Online Reinforcement Learning.

4.1 OFFLINE SUPERVISED LEARNING

The motivation behind training our student models in an offline supervised setting was primarily to test the usefulness of the ability to ask questions. The dialogues are generated as described in

the previous section, and the bot’s role is generated with a fixed policy. We chose a policy where answers to the teacher’s questions are correct answers 50% of the time, and incorrect otherwise, to add a degree of realism. Similarly, in tasks where questions can be irrelevant they are only asked correctly 50% of the time.⁶

The offline setting explores different combinations of training and testing scenarios, which mimic different situations in the real world. The aim is to understand when and how observing interactions between two agents can help the bot improve its performance for different tasks. As a result we construct training and test sets in three ways across all tasks, resulting in 9 different scenarios per task, each of which correspond to a real world scenario.

The three training sets we generated are referred to as **TrainQA, TrainAQ, and TrainMix**. TrainQA follows the QA setting discussed in the previous section: the bot never asks questions and only tries to immediately answer. TrainAQ follows the AQ setting: the student, before answering, first always asks a question in response to the teacher’s original question. TrainMix is a combination of the two where 50% of time the student asks a question and 50% does not.

The three test sets we generated are referred to as **TestQA, TestAQ, and TestModelAQ**. TestQA and TestAQ are generated similarly to TrainQA and TrainAQ, but using a perfect fixed policy (rather than 50% correct) for evaluation purposes. The setting of the TestModelAQ is different from the rest. Here while deciding on the question to be asked by the student, instead of using a fixed perfect policy which always results in a relevant question, we use a learned model to generate the question. The learned model is trained using either the TrainAQ or TrainMix set, depending on the training scenario. The teacher will reply to the question, giving positive feedback if the student’s question is correct and negative feedback otherwise. The student will then give the final answer. The difference between TestModelAQ and TestAQ only exists in the *Question Verification* and *Knowledge Verification* tasks; in other tasks there is only one way to ask the question and TestModelAQ and TestAQ are identical.

To summarize, for every task listed in Section 3 we train one model for each of the three training sets (TrainQA, TrainAQ, TrainMix) and test each of these models on the three test sets (TestQA, TestAQ, and TestModelAQ), resulting in 9 combinations. For the purpose of notation the train/test combination is denoted by “TrainSetting+TestSetting”. For example, *TrainAQ+TestQA* denotes a model which is trained using the TrainAQ dataset and tested on TestQA dataset. Each combination has a real world interpretation. For instance, *TrainAQ+TestQA* would refer to a scenario where a student can ask the teacher questions during learning but cannot to do so while taking an exam. Similarly, *TrainQA+TestQA* describes a stoic teacher that never answers a student’s question at either learning or examination time. The setting *TrainQA+TestAQ* corresponds to the case where a lazy student never asks question at learning time but gets anxious during the examination and always asks a question.

4.2 ONLINE REINFORCEMENT LEARNING (RL)

We also explored scenarios where the student learns the ability to decide when to ask a question and what to ask. In other words, the student learns how to learn.

Although it is in the interest of the student to ask questions at every step of the conversation, since the response to its question will contain extra information, we don’t want our model to learn this behavior. Each time a human student asks a question, there’s a cost associated with that action. This cost is a reflection of the patience of the teacher, or more generally of the users interacting with the bot in the wild: users won’t find the bot engaging if it always asks clarification questions. The student should thus be judicious about asking questions and learn when and what to ask. For instance, if the student is confident about the answer, there is no need for it to ask. Or, if the teacher’s question is so hard that clarification is unlikely to help enough to get the answer right, then it should also refrain from asking.

We now discuss how we model this problem under the Reinforcement Learning framework. The bot is presented with KB facts (some facts might be missing depending on the task) and a question. It needs to decide whether to ask a question or not at this point. The decision whether to ask is

⁶This only makes sense in tasks like Question or Knowledge Verification. In tasks where the question is static such as ‘What do you mean?’ there is no way to ask an irrelevant question, and we do not use this policy.

made by a binary policy $P_{RLQuestion}$. If the student chooses to ask a question, it will be penalized by $cost_{AQ}$. We explored different values of $cost_{AQ}$ ranging from $[0, 2]$, which we consider as modeling the patience of the teacher. The goal of this setting is to find the best policy for asking/not-asking questions which would lead to the highest cumulative reward. The teacher will appropriately reply if the student asks a question. The student will eventually give an answer to the teacher’s initial question at the end using the policy $P_{RLAnswer}$, regardless of whether it had asked a question. The student will get a reward of $+1$ if its final answer is correct and -1 otherwise. Note that the student can ask at most one question and that the type of question is always specified by the task under consideration. The final reward the student gets is the cumulative reward over the current dialogue episode. In particular the reward structure we propose is the following:

	Asking Question	Not asking Question
Final Answer Correct	$1 - cost_{AQ}$	1
Final Answer Incorrect	$-1 - cost_{AQ}$	-1

Table 1: Reward structure for the Reinforcement Learning setting.

For each of the tasks described in Section 3, we consider three different RL scenarios.

Good-Student: The student will be presented with all relevant KB facts. There are no misspellings or unknown words in the teacher’s question. This represents a knowledgeable student in the real world that knows as much as it needs to know (e.g., a large knowledge base, large vocabulary). This setting is identical across all missing entity tasks (5 - 9).

Poor-Student: The KB facts or the questions presented to the student are flawed depending on each task. For example, for the *Question Clarification* tasks, the student does not understand the question due to spelling mistakes. For the *Missing Question Entity* task the entity that the teacher asks about is unknown by the student and all facts containing the entity will be hidden from the student. This setting is similar to a student that is underprepared for the tasks.

Medium-Student: The combination of the previous two settings where for 50% of the questions, the student has access to the full KB and there are no new words or phrases or entities in the question, and 50% of the time the question and KB are taken from the *Poor-Student* setting.

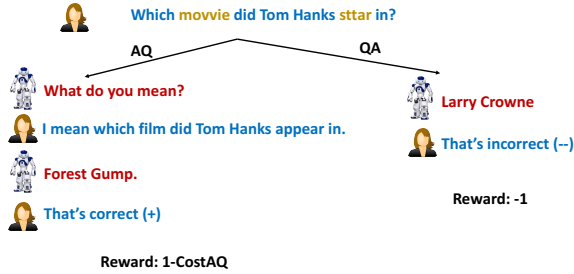


Figure 4: An illustration of the *poor-student* setting for RL Task 1 (Question Paraphrase).

5 MODELS

For both offline supervised and online RL settings, we use the *End-to-End Memory Network* model (MemN2N) (Sukhbaatar et al., 2015) as a backbone. The model takes as input the last utterance of the dialogue history (the question from the teacher) as well as a set of memory contexts including short-term memories (the dialogue history between the bot and the teacher) and long-term memories (the knowledge base facts that the bot has access to), and outputs a label. We refer readers to the Appendix for more details about MemN2N.

Offline Supervised Settings: The first learning strategy we adopt is the reward-based imitation strategy (denoted *vanilla-MemN2N*) described in (Weston, 2016), where at training time, the model maximizes the log likelihood probability of the correct answers the student gave (examples with incorrect final answers are discarded). Candidate answers are words that appear in the memories, which means the bot can only predict the entities that it has seen or known before.

We also use a variation of MemN2N called “context MemN2N” (*Cont-MemN2N* for short) where we replace each word’s embedding with the average of its embedding (random for unseen words) and the embeddings of the other words that appear around it. We use both the preceeding and following words as context and the number of context words is a hyperparameter selected on the dev set.

An issue with both *vanilla-MemN2N* and *Cont-MemN2N* is that the model only makes use of the bot’s answers as signals and ignores the teacher’s feedback. We thus propose to use a model that jointly predicts the bot’s answers and the teacher’s feedback (denoted as *TrainQA (+FP)*). The bot’s answers are predicted using a *vanilla-MemN2N* and the teacher’s feedback is predicted using the *Forward Prediction* (FP) model as described in (Weston, 2016). We refer the readers to the Appendix for the FP model details. At training time, the models learn to jointly predict the teacher’s feedback and the answers with positive reward. At test time, the model will only predict the bot’s answer.

For the *TestModelAQ* setting described in Section 4, the model needs to decide the question to ask. Again, we use *vanilla-MemN2N* that takes as input the question and contexts, and outputs the question the bot will ask.

Online RL Settings: A binary *vanilla-MemN2N* (denoted as $P_{RL}(Question)$) is used to decide whether the bot should or should not ask a question, with the teacher replying if the bot does ask something. A second MemN2N is then used to decide the bot’s answer, denoted as $P_{RL}(Answer)$. $P_{RL}(Answer)$ for *QA* and *AQ* are two separate models, which means the bot will use different models for final-answer prediction depending on whether it chooses to ask a question or not.⁷

We use the REINFORCE algorithm (Williams, 1992) to update $P_{RL}(Question)$ and $P_{RL}(Answer)$. For each dialogue, the bot takes two sequential actions (a_1, a_2) : to ask or not to ask a question (denoted as a_1); and guessing the final answer (denoted as a_2). Let $r(a_1, a_2)$ denote the cumulative reward for the dialogue episode, computed using Table 1. The gradient to update the policy is given by:

$$\begin{aligned} p(a_1, a_2) &= P_{RL}(Question)(a_1) \cdot P_{RL}(answer)(a_2) \\ \nabla J(\theta) &\approx \nabla \log p(a_1, a_2)[r(a_1, a_2) - b] \end{aligned} \quad (1)$$

where b is the baseline value, which is estimated using another MemN2N model that takes as input the query x and memory C , and outputs a scalar b denoting the estimation of the future reward. The baseline model is trained by minimizing the mean squared loss between the estimated reward b and actual cumulative reward r , $\|r - b\|^2$. We refer the readers to (Ranzato et al., 2015; Zaremba & Sutskever, 2015) for more details. The baseline estimator model is independent from the policy models and the error is not backpropagated back to them.

In practice, we find the following training strategy yields better results: first train only $P_{RL}(answer)$, updating gradients only for the policy that predicts the final answer. After the bot’s final-answer policy is sufficiently learned, train both policies in parallel⁸. This has a real-world analogy where the bot first learns the basics of the task, and then learns to improve its performance via a question-asking policy tailored to the user’s patience (represented by $cost_{AQ}$) and its own ability to answer questions.

6 EXPERIMENTS

Offline Results: Offline results are presented in Tables 2, 3 and 4. Table 2 presents results for the *vanilla-MemN2N* and *Forward Prediction* models. Table 3 presents results for *Cont-MemN2N*, which is better at handling unknown words. We repeat each experiment for 10 times and report the best result. Finally, Table 4 presents results for the test scenario where the bot itself chooses when to ask questions. Observations can be summarized as follows:

- Asking questions helps at test time, which is intuitive since it provides additional evidence:

⁷An alternative is to train one single model for final answer prediction in both *AQ* and *QA* cases, similar to the *TrainMix* setting in the supervised learning setting. But we find training *AQ* and *QA* separately for the final answer prediction yields a little better result than the single model setting.

⁸We implement this by running 16 epochs in total, updating only the model’s policy for final answers in the first 8 epochs while updating both policies during the second 8 epochs. We pick the model that achieves the best reward on the dev set during the final 8 epochs. Due to relatively large variance for RL models, we repeat each task 5 times and keep the best model on each task.

Question Clarification					Knowledge Operation			
Train \ Test	Task 1: Q. Paraphrase		Task 2: Q. Verification		Task 3: Ask For Relevant K.		Task 4: K. Verification	
	TestQA	TestAQ	TestQA	TestAQ	TestQA	TestAQ	TestQA	TestAQ
TrainQA	0.338	0.284	0.340	0.271	0.462	0.344	0.482	0.322
TrainAQ	0.213	0.450	0.225	0.373	0.187	0.632	0.283	0.540
TrainAQ(+FP)	0.288	0.464	0.146	0.320	0.342	0.631	0.311	0.524
TrainMix	0.326	0.373	0.329	0.326	0.442	0.558	0.476	0.491

Knowledge Acquisition										
Train \ Test	TestQA	TestAQ	TestQA	TestAQ	TestQA	TestAQ	TestQA	TestAQ	TestQA	TestAQ
	Task 5: Q. Entity		Task 6: Answer Entity		Task 7: Relation Entity		Task 8: Triple		Task 9: Everything	
TrainQA (vanila)	< 0.01	0.223	<0.01	<0.01	0.109	0.129	0.201	0.259	<0.01	<0.01
TrainAQ (vanila)	< 0.01	0.660	<0.01	<0.01	0.082	0.156	0.124	0.664	<0.01	<0.01
TrainAQ(+FP)	< 0.01	0.742	< 0.01	< 0.01	0.085	0.188	0.064	0.702	<0.01	<0.01
Mix (vanila)	<0.01	0.630	<0.01	<0.01	0.070	0.152	0.180	0.572	<0.01	<0.01

Table 2: Results for offline settings using memory networks.

Question Clarification					Knowledge Operation			
Train \ Test	Task 1: Q. Paraphrase		Task 2: Q. Verification		Task 3: Ask For Relevant K.		Task 4: K. Verification	
	TestQA	TestAQ	TestQA	TestAQ	TestQA	TestAQ	TestQA	TestAQ
TrainQA (Context)	0.754	0.726	0.742	0.684	0.883	0.947	0.888	0.959
TrainAQ (Context)	0.640	0.889	0.643	0.807	0.716	0.985	0.852	0.987
TrainMix (Context)	0.751	0.846	0.740	0.789	0.870	0.985	0.875	0.985

Knowledge Acquisition										
Train \ Test	Task 5: Q. Entity		Task 6: Answer Entity		Task 7: Relation Entity		Task 8: Triple		Task 9: Everything	
	TestQA	TestAQ	TestQA	TestAQ	TestQA	TestAQ	TestQA	TestAQ	TestQA	TestAQ
TrainQA (Context)	<0.01	0.224	<0.01	0.120	0.241	0.301	0.339	0.251	<0.01	0.058
TrainAQ (Context)	<0.01	0.639	<0.01	0.885	0.143	0.893	0.154	0.884	<0.01	0.908
TrainMix (Context)	<0.01	0.632	<0.01	0.852	0.216	0.898	0.298	0.886	<0.01	0.903

Table 3: Results for Cont-MemN2N on different tasks.

	Question Clarification	Knowledge Acquisition
	Task 2: Q. Verification	Task 4: K. Verification
	TestModelAQ	TestModelAQ
TrainAQ	0.382	0.480
TrainAQ(+FP)	0.344	0.501
TrainMix	0.352	0.469

Table 4: Results for TestModelAQ settings.

- *TrainAQ+TestAQ* (questions can be asked at both training and test time) performs the best across all the settings.
- *TrainQA+TestAQ* (questions can be asked at training time but not at test time) performs worse than *TrainQA+TestQA* (questions can be asked at neither training nor test time) in tasks *Question Clarification* and *Knowledge Operation* due to the discrepancy between training and testing.
- *TrainQA+TestAQ* performs better than *TrainQA+TestQA* on all *Knowledge Acquisition* tasks, the only exception being the *Cont-MemN2N* model on the *Missing Triple* setting. The explanation is that for most tasks in *Knowledge Acquisition*, the learner has no chance of giving the correct answer without asking questions. The benefit from asking is thus large enough to compensate for the negative effect introduced by data discrepancy between training and test time.
- *TrainMix* offers flexibility in bridging the gap between datasets generated using QA and AQ, very slightly underperforming *TrainAQ+TestAQ*, but gives competitive results on both *TestQA* and *TestAQ* in the *Question Clarification* and *Knowledge Operations* tasks.
- *TrainAQ+TestQA* (allowing questions at training time but forbid questions at test time) performs the worst, even worse than *TrainQA+TestQA*. This has a real-world analogy where a student becomes dependent on the teacher answering their questions, later struggling to answer the test questions without help.
- In the *Missing Question Entity* task (the student does not know about the question entity), the *Missing Answer Entity* task (the student does not know about the answer entity), and *Missing Everything* task, the bot achieves accuracy less than 0.01 if not asking questions at test time (i.e., *TestQA*).

- The performance of *TestModelAQ*, where the bot relies on its model to ask questions at test time (and thus can ask irrelevant questions) is worse than asking the correct question at test time (*TestAQ*) but better than not asking questions (*TestQA*).

- *Cont-MemN2N* significantly outperforms *vanilla-MemN2N*. One explanation is that considering context provides significant evidence distinguishing correct answers from candidates in the dialogue history, especially in cases where the model encounters unfamiliar words.

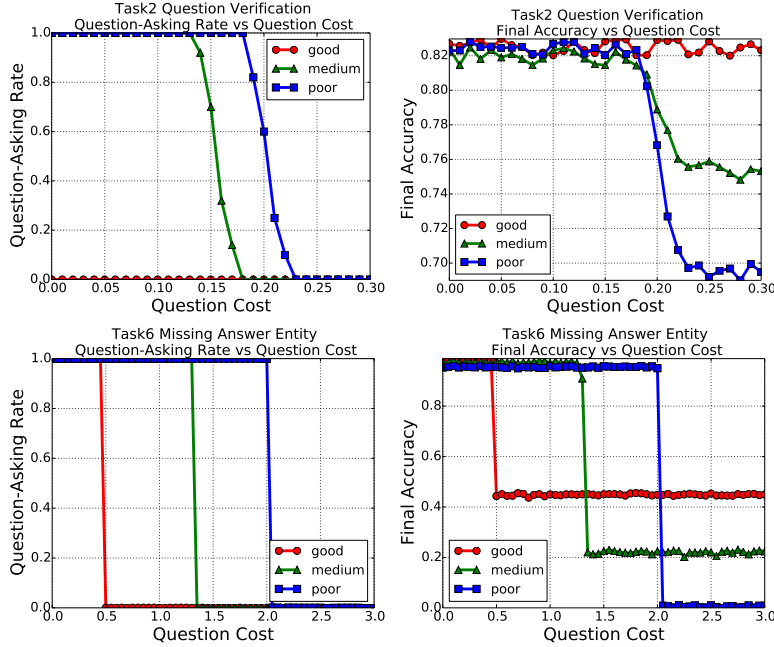


Figure 5: Results of online learning for Task 2 and Task 6

RL Results For the RL settings, we present results for Task 2 (*Question Verification*) and Task 6 (*Missing Answer Entities*) in Figure 5. Task 2 represents scenarios where different types of student have different abilities to correctly answer questions (e.g., a poor student can still sometimes give correct answers even when they do not fully understand the question). Task 6 represents tasks where a poor learner who lacks the knowledge necessary to answer the question can hardly give a correct answer. All types of students including the good student will theoretically benefit from asking questions (asking for the correct answer) in Task 6. We show the percentage of question-asking versus the cost of AQ on the test set and the accuracy of question-answering on the test set vs the cost of AQ. Our main findings were:

- A good student does not need to ask questions in Task 2 (*Question Verification*), because they already understand the question. The student will raise questions asking for the correct answer when cost is low for Task 6 (*Missing Answer Entities*).
- A poor student always asks questions when the cost is low. As the cost increases, the frequency of question-asking declines.
- As the AQ cost increases gradually, good students will stop asking questions earlier than the medium and poor students. The explanation is intuitive: poor students benefit more from asking questions than good students, so they continue asking even with higher penalties.
- As the probability of question-asking declines, the accuracy for poor and medium students drops. Good students are more resilient to not asking questions.

7 CONCLUSIONS

In this paper, we explored how an intelligent agent can benefit from interacting with users by asking questions. We developed tasks where interaction via asking questions is desired. We explore both online and offline settings that mimic different real world situations and show that in most cases, teaching a bot to interact with humans facilitates language understanding, and consequently leads to better question answering ability.

REFERENCES

- Mohammad Amin Bassiri. Interactional feedback and the impact of attitude and motivation on noticing l2 form. *English Language and Literature Studies*, 1(2):61, 2011.
- Antoine Bordes and Jason Weston. Learning end-to-end goal-oriented dialog. *arXiv preprint arXiv:1605.07683*, 2016.
- Antoine Bordes, Nicolas Usunier, Sumit Chopra, and Jason Weston. Large-scale simple question answering with memory networks. *arXiv preprint arXiv:1506.02075*, 2015.
- Jesse Dodge, Andreea Gane, Xiang Zhang, Antoine Bordes, Sumit Chopra, Alexander Miller, Arthur Szlam, and Jason Weston. Evaluating prerequisite qualities for learning end-to-end dialog systems. *arXiv preprint arXiv:1511.06931*, 2015.
- Richard Higgins, Peter Hartley, and Alan Skelton. The conscientious consumer: Reconsidering the role of assessment feedback in student learning. *Studies in higher education*, 27(1):53–64, 2002.
- Andrew S Latham. Learning through feedback. *Educational Leadership*, 54(8):86–87, 1997.
- Jiwei Li, Michel Galley, Chris Brockett, Jianfeng Gao, and Bill Dolan. A diversity-promoting objective function for neural conversation models. *arXiv preprint arXiv:1510.03055*, 2015.
- Alexander Miller, Adam Fisch, Jesse Dodge, Amir-Hossein Karimi, Antoine Bordes, and Jason Weston. Key-value memory networks for directly reading documents. *arXiv preprint arXiv:1606.03126*, 2016.
- Marc’Aurelio Ranzato, Sumit Chopra, Michael Auli, and Wojciech Zaremba. Sequence level training with recurrent neural networks. *arXiv preprint arXiv:1511.06732*, 2015.
- Alessandro Sordani, Michel Galley, Michael Auli, Chris Brockett, Yangfeng Ji, Margaret Mitchell, Jian-Yun Nie, Jianfeng Gao, and Bill Dolan. A neural network approach to context-sensitive generation of conversational responses. *arXiv preprint arXiv:1506.06714*, 2015.
- Pei-Hao Su, Milica Gasic, Nikola Mrksic, Lina Rojas-Barahona, Stefan Ultes, David Vandyke, Tsung-Hsien Wen, and Steve Young. Continuously learning neural dialogue management. *arXiv preprint arXiv:1606.02689*, 2016.
- Sainbayar Sukhbaatar, Jason Weston, Rob Fergus, et al. End-to-end memory networks. In *Advances in neural information processing systems*, pp. 2440–2448, 2015.
- Oriol Vinyals and Quoc Le. A neural conversational model. *arXiv preprint arXiv:1506.05869*, 2015.
- Sida I Wang, Percy Liang, and Christopher D Manning. Learning language games through interaction. *arXiv preprint arXiv:1606.02447*, 2016.
- Tsung-Hsien Wen, Milica Gasic, Nikola Mrksic, Lina M Rojas-Barahona, Pei-Hao Su, Stefan Ultes, David Vandyke, and Steve Young. A network-based end-to-end trainable task-oriented dialogue system. *arXiv preprint arXiv:1604.04562*, 2016.
- Margaret G Werts, Mark Wolery, Ariane Holcombe, and David L Gast. Instructive feedback: Review of parameters and effects. *Journal of Behavioral Education*, 5(1):55–75, 1995.
- Jason Weston. Dialog-based language learning. *arXiv preprint arXiv:1604.06045*, 2016.

Jason Weston, Antoine Bordes, Sumit Chopra, Alexander M Rush, Bart van Merriënboer, Armand Joulin, and Tomas Mikolov. Towards ai-complete question answering: A set of prerequisite toy tasks. *arXiv preprint arXiv:1502.05698*, 2015.

Ronald J Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine learning*, 8(3-4):229–256, 1992.

Terry Winograd. Understanding natural language. *Cognitive psychology*, 3(1):1–191, 1972.

Ludwig Wittgenstein. *Philosophical investigations*. John Wiley & Sons, 2010.

Wojciech Zaremba and Ilya Sutskever. Reinforcement learning neural turing machines. *arXiv preprint arXiv:1505.00521*, 362, 2015.

Appendix

End-to-End Memory Networks The input to an end-to-end memory network model (MemN2N) is the last utterance of the dialogue history x as well as a set of memories (context) ($C=c_1, c_2, \dots, c_N$). Memory C encodes both short-term memory, e.g, dialogue histories between the bot and the teacher and long-term memories, e.g., the knowledgebase facts that the bot has access to. Given the input x and C , the goal is to produce an output/label a .

In the first step, the query x is transformed to a vector representation u_0 by summing up its constituent word embeddings: $u_0 = Ax$. The input x is a bag-of-words vector and A is the $d \times V$ word embedding matrix where d denotes the vector dimensionality and V denotes the vocabulary size. Each memory c_i is similarly transformed to vector m_i . The model will read information from the memory by linking input representation q with memory vectors m_i using softmax weights:

$$o_1 = \sum_i p_i^1 m_i \quad p_i^1 = \text{softmax}(u_0^T m_i) \quad (2)$$

The goal is to select memories relevant to the last utterance x , i.e., the memories with large values of p_i^1 . The queried memory vector o_1 is the weighted sum of memory vectors. The queried memory vector o_1 will be added on top of original input, $u_1 = o_1 + u_0$. u_1 is then used to query the memory vector. Such a process is repeated by querying the memory N times (so called “hops”). N is set to three in all experiments in this paper.

In the end, u_N is input to a softmax function for the final prediction:

$$a = \text{softmax}(u_N^T y_1, u_N^T y_2, \dots, u_N^T y_L) \quad (3)$$

where L denotes the number of candidate answers and y denotes the representation of the answer. If the answer is a word, y is the corresponding word embedding. If the answer is a sentence, y is the embedding for the sentence achieved in the same way as we obtain embeddings for query x and memory c .

Reward Based Imitation (RBI) and Forward Prediction (FP) *RBI* and *FP* are two dialogue learning strategies proposed in (Weston, 2016) by harnessing different types of dialogue signals. *RBI* handles the case where the reward or the correctness of a bot’s answer is explicitly given (for example, +1 if the bot’s answer is correct and 0 otherwise). The model is directly trained to predict the correct answers (with label 1) at training time, which can be done using End-to-End Memory Networks (*MemN2N*) (Sukhbaatar et al., 2015) that map a dialogue input to a prediction.

FP handles the situation where a real-valued reward for a bot’s answer is not available, meaning that there is no +1 or 0 labels paired with a student’s utterance. However, the teacher will give a response to the bot’s answer, taking the form of a dialogue utterance. More formally, suppose that x denotes the teacher’s question and $C=c_1, c_2, \dots, c_N$ denotes the dialogue history. In our *AQ* settings, the bot will ask a question a regarding the teacher’s question, denoted as $a \in \mathbb{A}$, where \mathbb{A} denotes the student’s question pool. The teacher will provide an utterance in response to the student question a . In *FP*, the model first maps the teacher’s initial question x and dialogue history C to vector representation u using a memory network with multiple hops. Then the model will perform another hop of attention over all possible student’s questions in \mathbb{A} , with an additional part that incorporates the information of which candidate (i.e., a) was actually selected in the dialogue:

$$p_{\hat{a}} = \text{softmax}(u^T y_{\hat{a}}) \quad o = \sum_{\hat{a} \in \mathbb{A}} p_{\hat{a}} (y_{\hat{a}} + \beta \cdot \mathbf{1}[\hat{a} = a]) \quad (4)$$

where $y_{\hat{a}}$ denotes the vector representation for the student’s question candidate \hat{a} . β is a d -dimensional vector to signify the actual action a that the student chooses. For tasks where the student only has one way to ask questions (e.g., “what do you mean”), there is no need to perform hops of attention over candidates since the cardinality of \mathbb{A} is just 1. We thus directly assign a probability of 1 to the student’s question, making o the sum of vector representation of y_a and β .

o is then combined with u to predict the teacher’s feedback t using a softmax:

$$u_1 = o + u \quad t = \text{softmax}(u_1^T x_{r_1}, u_1^T x_{r_2}, \dots, u_1^T x_{r_N}) \quad (5)$$

where x_{r_i} denotes the embedding for the i^{th} response.

Dialogue Simulator In this section we further detail the simulator and the datasets we generated in order to realize the various scenarios discussed in Section 3. We focused on the problem of movieQA where we adapted the WikiMovies dataset proposed in Weston et al. (2015). The dataset consists of roughly 100k questions with over 75k entities from the open movie dataset (OMDb).

Each dialogue generated by the simulator takes place between a student and a teacher. The simulator samples a random question from the WikiMovies dataset and fetches the set of all KB facts relevant to the chosen question. This question is assumed to be the one the teacher asks its student, and is referred to as the “original” question. The student is first presented with the relevant KB facts followed by the original question. Providing the KB facts to the student allows us to control the exact knowledge the student is given access to while answering the questions. At this point, depending on the task at hand and the student’s ability to answer, the student might choose to directly answer it or ask a “followup” question. The nature of the followup question will depend on the scenario under consideration. If the student answers the question, it gets a response from the teacher about its correctness and the conversation ends. However if the student poses a followup question, the teacher gives an appropriate response, which should give additional information to the student to answer the original question. In order to make things more complicated, the simulator pads the conversation with several unrelated student-teacher question-answer pairs. These question-answer pairs can be viewed as distractions and are used to test the student’s ability to remember the additional knowledge provided by the teacher after it was queried. For each dialogue, the simulator incorporates 5 such pairs (10 sentences). We refer to these pairs as *conversational histories*.

For the *QA* setting (see Section 3), the dialogues generated by the simulator are such that the student never asks a clarification question. Instead, it simply responds to the original question, even if it is wrong. For the dialogs in the *AQ* setting, the student *always* asks a clarification question. The nature of the question asked is dependent on the scenario (whether it is *Question Clarification*, *Knowledge Operation*, or *Knowledge Acquisition*) under consideration. In order to simulate the case where the student sometimes chooses to directly answer the original question and at other times chooses to ask question, we created training datasets, which were a combination of *QA* and *AQ* (called “Mixed”). For all these cases, the student needs to give an answer to the teacher’s original question at the end.